

CAREER PROGRAM · AI ENGINEERING

AI Engineering

Build the AI systems companies are actually hiring for: retrieval that stays grounded, agents that use tools safely, and MCP servers that connect them to real systems. A six-month, project-first program whose syllabus is rebuilt every month from real AI Engineer interviews.

50/day

Real interviews monitored by our AI

~90%

Of interview questions come from our material

98%

Of candidates who attend interviews get placed

1 year

Unlimited live batches & recordings

6 MONTHS · LIVE ONLINE · BENGALURU · 2026 EDITION

EVERY PROMISE IN WRITING · EVERY CALL RECORDED & AI-MONITORED.

Most institutes teach a syllabus. We prepare you for the interview.

The problem with every other course: fixed syllabi written years ago, packed with topics no interviewer asks about — while the questions that decide your offer change every quarter. You spend months studying, then walk into interviews prepared for the wrong exam.

THE SYLLABUS ENGINE — REBUILT MONTHLY

This syllabus was not written. It was reverse-engineered.

STEP 01 — LISTEN

Our Claude-based AI monitors up to 50 real & mock interviews every single day.

STEP 02 — EXTRACT

It collects the actual questions companies are asking for this role, this month.

STEP 03 — REBUILD

The syllabus is rebuilt around real demand — you study only what's actually asked.

50/day

Real interviews monitored by our AI

~90%

Of interview questions come from our material

98%

Of candidates who attend interviews get placed

1 year

Unlimited live batches & recordings

*Based on BrowseJobs internal data. Historical figures — not a promise of individual outcome. Hiring depends on the live market and your performance.

Why AI Engineering is the fastest-moving track we run

Every company that adopted an LLM last year now needs someone who can make it reliable — grounded in their data, safe to give tools to, and cheap enough to run. That work has become its own discipline, and the syllabus for it changes faster than any other track we teach.

CAREER GROWTH FLOW

AI Engineer

Build RAG and agent features on top of existing models

Senior AI Engineer

Own evaluation, guardrails and cost for AI in production

AI Platform Engineer

Build the shared serving, retrieval and tooling layer

Lead / AI Architect

Set the architecture and the safety bar across teams

Built for your situation. Judged on your outcomes.

SWITCHING DOMAINS?

The real, industry-grade projects you build here become the centerpiece of your CV — hands-on work in the new domain, fully genuine, fully BGV-safe.

CAREER GAP?

A structured internship with us, built around projects you actually deliver, fills your timeline with documented work. A gap is only a gap if it's empty.

FRESHER?

Earn a free internship certificate with BrowseJobs.ai — the equivalent of six months of hands-on experience before your first interview.

WHAT YOU'LL WALK OUT ABLE TO DO

- ✓ Ship a retrieval system that answers from your data with citations — and refuses when it can't
- ✓ Build agents that call tools, keep memory and hand off work, without going off the rails
- ✓ Write and publish MCP servers that connect a model to real systems safely
- ✓ Prove a non-deterministic feature works: eval sets, tracing, regression tests and cost budgets
- ✓ Answer the RAG, agent-design and guardrail questions AI Engineer interviews lead with

ROLES THIS PREPARES YOU FOR

AI Engineer

LLM Application Engineer

Agentic AI Engineer

Applied AI Engineer

AI Platform Engineer

PROGRAM FORMAT

Duration	6 months
Format	Live online + recordings
Access	1 year unlimited
Projects	5 CV-ready

“We qualify you on data points, not emotions. Your mock-interview data decides when you're ready — and then the market meets you.”

— DR. KRISH BHARGAV, FOUNDER & MENTOR

The live syllabus. Rebuilt monthly from real interviews.

12 modules. Zero filler. Every topic below appeared in a monitored interview for this role — that is the only reason it is here.

MODULE 01

Python & Async for AI Engineering

The language floor every AI Engineering interview assumes.

- Core Python, typing, dataclasses, decorators
- asyncio and concurrent API calls — the difference between a demo and a service
- Pydantic models and schema validation
- Packaging, environments and dependency management with uv
- Git, pull requests and code review

MODULE 02

LLM Foundations & the Provider Layer

Understanding the model you are engineering around.

- Tokens, context windows, temperature, sampling and why output varies
- Comparing providers: capability, latency, price, rate limits
- One interface, many models — a provider-agnostic gateway you can swap
- Running models locally with Ollama, and when local is the right answer
- Streaming, timeouts, retries and graceful degradation

MODULE 03

Prompting, Context Engineering & Structured Outputs

The term that replaced 'prompt engineering' in this year's job descriptions.

- System prompts, few-shot patterns and instruction design
- Context engineering: what to put in the window, and what to leave out
- Structured outputs with JSON schema and Pydantic — parsing you can trust
- Chunking long inputs, summarisation chains and context compression
- Caching prompts and cutting token cost without losing quality

MODULE 04

RAG — Retrieval-Augmented Generation

The single most-asked topic in AI Engineer interviews.

- Embeddings, similarity and what a vector actually encodes
- Chunking strategies and why naive splitting breaks answers
- Vector stores: pgvector, Chroma, Qdrant — and choosing between them
- Hybrid search, metadata filtering and reranking
- Citations, groundedness and teaching a system to say 'I don't know'
- Measuring retrieval: recall, precision and failure analysis

MODULE 05

Agent Fundamentals & Design Patterns

What an agent is, and — more usefully — when not to build one.

- The agent loop: reason, act, observe, repeat
- Agentic design patterns: reflection, planning, tool use, multi-agent
- Tool calling: definitions, schemas, parallel calls, error recovery
- Memory: short-term, long-term and retrieval-backed
- Agentic AI risks — and the failure modes interviewers ask you to name

MODULE 06

Building Agents with the Major Frameworks

Depth in the frameworks JDs name, not a tour of every library.

- LangChain core: models, tools, structured output binding
- LangGraph: state, nodes, edges, checkpointing and durable execution
- Human-in-the-loop interrupts and approval gates
- The OpenAI Agents SDK: agents, handoffs, sessions
- CrewAI compared — roles, tasks and when a crew fits better
- Choosing a framework, and defending the choice in an interview

MODULE 07

Multi-Agent Systems & Orchestration

Where most candidates' understanding runs out.

- Orchestrator-worker, planner-executor and debate patterns
- Handoffs, delegation and shared state between agents
- The A2A protocol: agent-to-agent communication
- Sandboxed execution — running model-written code without losing the host
- Loop limits, budgets and stopping conditions

MODULE 08

MCP — Model Context Protocol

The integration layer that turned agents from demos into products.

- Hosts, clients and servers — the MCP architecture
- Tools, resources and prompts: the three things a server exposes
- Local (stdio) vs remote (HTTP/SSE) transport
- Connecting an agent to existing MCP servers and marketplaces
- Writing your own MCP server over a real system, with auth
- Agentic RAG and long-term memory through MCP

MODULE 09

Guardrails, Safety & Prompt-Injection Defence

The module that decides whether your system is allowed near production.

- Input and output guardrails, and where each belongs
- Prompt injection, indirect injection and the lethal trifecta
- Least-privilege tool access and read-only by default
- PII handling, redaction and data-residency constraints
- Human approval gates and audit trails for agent actions
- Refusal, fallback and failing safely

MODULE 10

Evaluation, Observability & LLMOps

How you prove a non-deterministic feature works — the question that ends most interviews.

- Building an eval set from real usage, not vibes
- LLM-as-a-judge: rubrics, bias, and validating the judge itself
- Regression testing prompts and models before you ship
- Tracing with LangSmith: spans, latency, token accounting
- Cost budgets, model routing and caching strategy
- Monitoring drift and quality in production

MODULE 11

Deploying AI Systems to Production

Deployed, monitored and affordable — or it isn't engineering.

- Serving agents behind FastAPI with streaming responses
- Containerising an AI service and managing model configuration
- Queues and background execution for long-running agent runs
- Rate limiting, quotas and per-user token budgets
- Building a UI with Gradio, and shipping it
- CI/CD, staged rollout and rollback for prompt and model changes

Interview Engineering

The module no other institute has — because no one else monitors 50 interviews a day.

- Live question bank from monitored AI Engineer interviews, refreshed monthly
- Weekly technical + HR mock interviews, AI-analysed
- System-design drills: design a RAG platform, design an agent runtime
- Communication & confidence coaching — scored, not optional

Real projects you'll build. These go on your CV.

Not toy exercises — industry-grade work you deliver, deploy and can walk an interviewer through line by line. Fully genuine, so every candidate clears background verification.

PROJECT 01

Grounded knowledge assistant

RAG over a real document corpus — hybrid search, reranking, citations, and a measured refusal rate.

Chunking & embeddings

Hybrid search + rerank

Cited answers

Retrieval eval set

PROJECT 02

Autonomous research analyst

A planner, searcher and writer working together to produce a sourced report from one question.

Multi-agent orchestration

Web search tools

Structured report output

Loop & cost limits

PROJECT 03

Your own MCP server

An MCP server over a real system, authenticated, connected to a host and driven by an agent.

Tools, resources, prompts

Local & remote transport

Auth & least privilege

Published & documented

PROJECT 04

Guardrailed agent with human-in-the-loop

An agent that takes real actions — behind input/output guardrails, approval gates and an audit trail.

Injection defence

Approval interrupts

Audit log

Safe failure paths

PROJECT 05

Eval harness & cost dashboard

The evidence layer: regression tests over prompts and models, traced, with token spend per feature.


Eval set & rubrics

LLM-as-a-judge

LangSmith tracing

Cost per request

THE STACK YOU'LL HAVE USED



COACH NOTE • EXPLORE LATER

The core program goes deep on the frameworks named in live job descriptions rather than touring every library. If a specific JD asks for one of these, here is your self-study map — and your mentor will point you at it when it matters.

- Google ADK — agents on Google's stack
- Pydantic AI — typed agents in pure Python
- Microsoft Agent Framework and Agno
- Mastra — building agents in TypeScript
- n8n — low-code agent workflows
- Fine-tuning and parameter-efficient adaptation (LoRA)

Your projects don't stop at the notebook.

Every CV-ready project in this program is deployed, not just coded. An interviewer can open it, and you can walk them through the pipeline that put it there — which is exactly the conversation that converts a project into an offer.

01

Your own cloud account

You build in a real AWS or Azure account with real IAM, real billing alarms and real limits — not a simulator.

02

Version-controlled from day one

Git and GitHub from the first commit, with branches, pull requests and review — the workflow your first team will expect.

03

Automated build & deploy

A pipeline takes each project from commit to running environment, so deployment is a repeatable step rather than a one-off.

04

Monitored and demoable

Dashboards, logs and alerts on top, plus a written architecture walkthrough you rehearse in mock interviews.

Everything included. The whole machine, from day one.

No add-on pricing for the things that get you hired. Every tool below ships with this track.

LEARN

PROVE

GET HIRED

AI Tutor — awake at 2am

Every doubt answered in seconds, in your context, from your syllabus. Students ask it hundreds of questions a course, because it never judges and never sleeps.

Live batches + recordings

Every class is live and instructor-led, then recorded and searchable by topic. Your access runs a full year, so a missed class is never a lost one.

Coding Lab

Browser-based labs with the real stack. You write, run and break things inside the LMS — no local setup standing between you and practice.

Content Hub & lesson notes

Per-lesson notes, slides and reference material, generated against the day's teaching plan and kept with the recording.

Voice Mock Lab — the interviewer that calls you

A voice AI rings your phone, runs a real technical round, and scores every answer on depth, communication and confidence. By the twelfth mock, the real thing feels like a rerun.

Weekly technical + HR mocks

Human-led mocks alongside the AI ones, AI-analysed afterwards, each returning a written gap report rather than a grade.

PRI & the mastery map

A Placement Readiness Index built from your mock data, attendance and assignments — with one clear next best action every day.

Live interview question bank

The questions pulled from monitored interviews for your role this month, organised round by round.

Quizzes, assignments & graded feedback

Scored work with trainer feedback, so readiness is a measurement rather than an opinion.

Job Radar

Live roles from LinkedIn and Naukri matched to your skill graph and refreshed daily, so the openings are already waiting when you are ready.

ATS resume, LinkedIn & GitHub

Your CV rebuilt to survive the filter, your LinkedIn tuned for recruiter search, your GitHub arranged around the projects you actually shipped.

Verifiable certificates

Every certificate carries a QR code that resolves to a public verification page. An employer checks it in one scan.

1:1 mentoring & support desk

A mentor for the career decisions and a ticketed support desk for everything else, both inside the LMS.

The placement engine. You don't apply into a black hole.

You get presented — through three channels, by our team, with your profile optimised first.

01 · **NAUKRI, DONE RIGHT**

Profile built by our team

Your CV professionally optimised and loaded on Naukri by our team — not left to you.

02 · **3,000+ HR NETWORK**

Broadcast to recruiters

Your profile shared across our network of recruiters and HR professionals actively hiring.

03 · **DIRECT CLIENTS**

Straight to hiring companies

Shared directly with our client companies hiring for this role right now.

CAREER SERVICES INCLUDED

SERVICE	WHAT YOU GET
ATS Resume Building	A resume that survives the filter and gets read by a human
LinkedIn & GitHub	Profile optimised for recruiter search; portfolio that showcases real work
Mock Interviews	Weekly technical + HR mocks, AI-analysed, with detailed feedback
Communication Training	A scored module — not an add-on. Delivery is what converts skill into offers
Interview Question Bank	Live bank drawn from monitored interviews for this role, this month
1:1 Mentoring	Personalised guidance at every step until you're hired

READINESS IS DECIDED ON DATA

We send you to interviews when your mock scores say you're ready — not when a sales target says so. That is precisely why our attend-to-place rate is what it is.

Fees that bet on you. Most of it only after you're placed.

YOUR PATH — EVERYTHING IS FREE UNTIL STEP 04

FREE • 01

Free counselling + written Career Analysis Report

Talk to us. You leave with a written report on where you stand and what to do — whether you join or not.

FREE • 02

Free live Masterclass

See the method live. How the syllabus is reverse-engineered from real interviews, and what you'd actually study.

FREE • 03

Free 7-hour Python Bootcamp

Sit a real teaching day before you decide. Seven hours of live Python — no card, no commitment.

STEP 04

Registration — ₹30,000

Only after all three free steps. EMI available: 3 × ₹10,000. Then full enrolment and LMS access.

Registration (payable only after the free masterclass + bootcamp)

₹30,000

EMI option for registration

3 × ₹10,000

Placement fee — due only after you accept an offer

First 3 months' CTC

Paid as

6 monthly EMIs from your new salary

Your ₹30,000 registration

Adjusted inside the placement fee

WORKED EXAMPLE • ₹12 LPA OFFER

3 months' CTC = ₹3,00,000 — ₹30,000 already paid = ₹2,70,000 in 6 EMIs — paid from a salary you didn't have before. We only earn properly when you do.

30-day money-back guarantee

Any reason, or no reason — full refund in your first 30 days of the program. No questions asked. Written into your agreement.

Don't trust us. Verify us.

Every institute will tell you their course is the best. So don't decide on anyone's words — including ours. Here is exactly how to check the market yourself, tonight, in fifteen minutes, for free.

CHECK 01 · 5 MIN

Count the demand

Open Naukri → search the role + your city → filter “Last 7 days” → count the postings. If hiring is real, you'll see it in the number.

CHECK 02 · 5 MIN

Read 5 job descriptions

Note the skills that repeat. Then compare them against any institute's syllabus — ours included. A syllabus that doesn't match live JDs is preparing you for the wrong exam.

CHECK 03 · 5 MIN

Pressure-test everyone

Ask every institute: will you put every promise in writing? Is your success rate defined clearly? Do you guarantee a job? Compare the answers.

What we promise — in writing

- ✓ Every commitment you get from us is in writing.
- ✓ Three full steps are free before you pay a rupee.
- ✓ The placement fee is due only after you accept an offer.
- ✓ 30-day money-back guarantee — any reason, full refund.
- ✓ Every counselling call is recorded & AI-monitored.

What we will never tell you

- ✗ “We guarantee you a job.” Nobody honestly can — the market decides.
- ✗ “We'll adjust your experience.” We never fabricate employment.
- ✗ A fixed salary promise.
- ✗ That hiring is certain, or that there's a shortcut.

QUESTIONS TO ASK ANY INSTITUTE — INCLUDING US

- Will you put every promise — fees, refund, placement process — in a written agreement?
- Is your success rate defined clearly? Ours: 98% of candidates who attend interviews, internal data.
- Do you guarantee a job? The only honest answer is no.
- Will anything on my CV be fabricated? With us: never — real projects and a real internship only.

Real stories. Real people. Real success.

★★★★★ DEVOPS

"Big thanks to Trainer Krish for being such a motivating and knowledgeable guide. Your way of teaching with real-life examples made everything so easy to understand. You're not just a trainer, but a true mentor!"

KR Sampritha

★★★★★ DATA ENGINEERING

"The course is super well-structured, covering everything from ETL processes to data storage tools. Instead of just theory, I got to work on actual data pipelines and industry-relevant case studies."

Anish Lakumarapu

★★★★★ DATA ENGINEERING

"I would like to sincerely thank BrowseJobs for their incredible support in helping me find a job. Krish Sir, you are doing a wonderful job in helping students who dream of working in the IT industry."

Saket Vaibhav

★★★★★ DEVOPS

"I recommend BrowseJobs to anyone looking to advance their career in IT. The course was comprehensive, well-structured, and tailored. My tutor was always available to clear any doubts."

Arbaz Khan

5,000+

Learners trained & empowered

4.9/5

Average rating, based on reviews

400+

Hiring partners connected

2013

London est. · India since 2020

*Based on BrowseJobs internal data. Historical figures – not a promise of individual outcome. Hiring depends on the live market and your performance.

Start with everything free.

Counselling + a written career report · the live masterclass · a 7-hour Python bootcamp. Then — and only then — decide.

Call or WhatsApp +91 86185 19825 · hello@browsejobs.ai · browsejobs.ai

BrowseJobs is a unit of IBrowseJobs Technologies Pvt Ltd · Fl 3, Channasandra Main Rd, Ambedkar Nagar, Whitefield, Bengaluru, Karnataka 560066 · Mon-Sat, 9:00 AM – 7:00 PM IST.